

HelixTriage — AI Use Case Requirements

System: HelixTriage · Reference: RAI-RD-20260415-helixtrriage · Owner: Enterprise Service Management · Version 1.2

System name	HelixTriage
Reference ID	RAI-RD-20260415-helixtrriage
Business owner	Enterprise Service Management
Lifecycle stage	Pilot — approved for limited production
Primary model	Azure OpenAI GPT-4o

1. Purpose

HelixTriage is an AI-assisted triage capability for the enterprise IT service desk. It reads each inbound support ticket, predicts the resolver group and priority, and drafts a suggested first response for an agent to approve or edit. The system does not close tickets autonomously; every prediction is presented to a human agent before it reaches the requester.

2. Business problem

Roughly 41% of tickets are currently routed to the wrong resolver group on first assignment. Each misroute adds an average of 6.4 hours to resolution time and consumes analyst capacity in manual reassignment. Volume has grown 23% year on year while service desk headcount has remained flat.

3. Scope

- In scope: English-language tickets from the web portal, email intake and the Teams bot.
- In scope: category prediction, priority prediction, suggested first response.
- Out of scope: autonomous resolution, ticket closure, entitlement or access decisions.
- Out of scope: tickets flagged as containing personal or clinical data at intake.

4. Functional requirements

ID	Requirement	Priority
FR-01	HelixTriage shall predict a resolver group for every in-scope ticket within 4 seconds of intake.	Must
FR-02	The system shall return a calibrated confidence score with every prediction.	Must
FR-03	Predictions below the confidence threshold shall be routed to manual triage unchanged.	Must
FR-04	The system shall draft a suggested first response for agent review.	Should
FR-05	An agent shall be able to override any prediction, and the override shall be recorded.	Must
FR-06	The system shall log model version and prompt version against every prediction.	Must

5. Non-functional requirements

Inference latency at the 95th percentile must remain under 4 seconds. HelixTriage must sustain 900 predictions per hour at peak. Availability target is 99.5% during business hours. All inference traffic must remain within the corporate network boundary; the inference endpoint `helixtrriage-inference.svc.internal.corp` is not exposed to the public internet.

6. Success metrics

- First-assignment accuracy above 85% on the held-out evaluation set.
- Misroute rate reduced from 41% to below 15% within the pilot cohort.
- Agent acceptance of suggested responses above 60%.
- No increase in reopened-ticket rate relative to the manual baseline.

7. Governance

HelixTriage is registered under reference RAI-RD-20260415-helixtrriage. A human agent reviews every prediction before it reaches a requester, and predictions are retained for 90 days for audit. Model and prompt versions are recorded per prediction so any decision can be reconstructed. Related documents: the solution design, the triage evaluation results, the labelled evaluation dataset, and the pilot rollout plan.